

¶ PRACTICAL 1: Implement dataset versioning using DVC

GOAL: Version a dataset file using DVC and a local remote. **PRE-REQUISITE:** Git must be installed.

1. Setup Project Folders

```
mkdir dvc-practical  
mkdir dvc-remote  
cd dvc-practical
```

2. Create and Activate Virtual Environment

```
python -m venv venv  
.\\venv\\Scripts\\activate
```

3. Install Libraries

```
pip install dvc
```

4. Initialize Git and DVC

```
git init  
dvc init
```

5. Configure a LOCAL remote

```
# Using '..' to go up one level from 'dvc-practical'  
dvc remote add -d local_remote ..\\dvc-remote
```

6. Create V1 of the dataset

```
echo "feature1,feature2,target" > data.csv  
echo "1,2,0" >> data.csv  
echo "3,4,1" >> data.csv
```

7. Track V1

```
dvc add data.csv  
git add data.csv.dvc .gitignore  
git commit -m "Add V1 of data.csv"
```

8. Push V1

```
dvc push -r local_remote
```

9. Create V2 of the dataset

```
echo "5,6,0" >> data.csv  
echo "7,8,1" >> data.csv
```

10. Track V2

```
dvc add data.csv  
git add data.csv.dvc  
git commit -m "Add V2 of data.csv"
```

11. Push V2

```
dvc push -r local_remote
```

12. Simulate restoring data

```
del data.csv
dvc pull -r local_remote
type data.csv
# (This will show V2, the latest version)
```

13. Switch back to V1

```
git checkout HEAD~1 data.csv.dvc
dvc checkout
type data.csv
# (This will show V1)
```

14. Switch back to V2 (latest)

```
git checkout main data.csv.dvc
dvc checkout
type data.csv
# (This will show V2)
```

¶ PRACTICAL 2: Track experiments using MLflow

GOAL: Train a model and log its parameters and metrics to the MLflow UI.

1. Setup Project Folder

```
mkdir mlflow-practical
cd mlflow-practical
```

2. Create and Activate Virtual Environment

```
python -m venv venv
.\venv\Scripts\activate
```

3. Install Libraries

```
pip install mlflow scikit-learn pandas
```

4. Create Python file: `create_data.py`

```
import pandas as pd
from sklearn.datasets import make_classification

X, y = make_classification(n_samples=100, n_features=10, n_informative=5, n_redundant=0, random_state=42)
df = pd.DataFrame(X, columns=[f'feature_{i}' for i in range(10)])
df['target'] = y

df.to_csv('dummy_data.csv', index=False)
print("Data created.")
```

5. Create Python file: `train.py`

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score
import mlflow
import mlflow.sklearn
import sys

# Set tracking URI to a local directory named 'mlruns'
mlflow.set_tracking_uri("file:./mlruns")

# Get C parameter from command line, default to 1.0
C_param = float(sys.argv[1]) if len(sys.argv) > 1 else 1.0

# Load data
df = pd.read_csv('dummy_data.csv')
X = df.drop('target', axis=1)
y = df['target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Start an MLflow run
with mlflow.start_run():
    print("Starting MLflow run...")
    # Train model
    model = LogisticRegression(C=C_param, max_iter=200, random_state=42)
    model.fit(X_train, y_train)

    # Make predictions
    preds = model.predict(X_test)

    # Calculate metrics
    accuracy = accuracy_score(y_test, preds)
    precision = precision_score(y_test, preds)

    # Log parameters
    mlflow.log_param("C", C_param)
    mlflow.log_param("model_type", "LogisticRegression")

    # Log metrics
    mlflow.log_metric("accuracy", accuracy)
    mlflow.log_metric("precision", precision)

    # Log the model
    mlflow.sklearn.log_model(model, "model")

    print(f"Run complete. Accuracy: {accuracy}")

```

6. Run setup and training

```

python create_data.py
python train.py 0.5
python train.py 1.0
python train.py 2.0

```

7. Launch MLflow UI

```
mlflow ui
```

Open your browser to <http://127.0.0.1:5000>

¶ PRACTICAL 3: Track experiments using Weights & Biases

GOAL: Train a model and log its parameters and metrics to the W&B dashboard. **PRE-REQUISITE:** You need a free W&B account (<https://wandb.ai/site>).

1. Setup Project Folder

```
mkdir wandb-practical  
cd wandb-practical
```

2. Create and Activate Virtual Environment

```
python -m venv venv  
.\\venv\\Scripts\\activate
```

3. Install Libraries

```
pip install wandb scikit-learn pandas
```

4. Login to W&B

```
wandb login
```

This will ask you to paste an API key from your W&B profile settings.

5. Create Python file: `create_data.py`

```
import pandas as pd  
from sklearn.datasets import make_classification  
  
X, y = make_classification(n_samples=100, n_features=10, n_informative=5, n_redundant=0, random_state=42)  
df = pd.DataFrame(X, columns=[f'feature_{i}' for i in range(10)])  
df['target'] = y  
  
df.to_csv('dummy_data.csv', index=False)  
print("Data created.")
```

6. Create Python file: `train_wandb.py`

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
import wandb

# Load data
df = pd.read_csv('dummy_data.csv')
X = df.drop('target', axis=1)
y = df['target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Configuration for the run
config = {
    "C": 1.0,
    "model_type": "LogisticRegression",
    "solver": "liblinear"
}

# 1. Initialize W&B run
wandb.init(project="practical-exam", config=config)

# 2. Train model
model = LogisticRegression(
    C=wandb.config.C,
    solver=wandb.config.solver,
    random_state=42
)
model.fit(X_train, y_train)

# Make predictions
preds = model.predict(X_test)
accuracy = accuracy_score(y_test, preds)

# 3. Log metrics
wandb.log({"accuracy": accuracy})

print(f"Run complete. Accuracy: {accuracy}")
wandb.finish()

```

7. Run setup and training

```

python create_data.py
python train_wandb.py

```

Check the terminal output for the W&B link to your project dashboard.

¶ PRACTICAL 4: Register models using MLflow Model Registry

GOAL: Log a model to the MLflow UI and register it by name. **NOTE:** This builds directly on Practical #2.

1. Setup Project Folder (Same as Practical #2)

```

mkdir mlflow-practical
cd mlflow-practical

```

2. Activate Environment (If not already active)

```
.\venv\Scripts\activate
```

3. Ensure `create_data.py` exists (from Practical #2)

4. Create Python file: `register_model.py`

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
import mlflow
import mlflow.sklearn

# --- This is the key model name for the registry ---
REGISTERED_MODEL_NAME = "PracticalExamModel"

mlflow.set_tracking_uri("file:./mlruns")

# Load data
df = pd.read_csv('dummy_data.csv')
X = df.drop('target', axis=1)
y = df['target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Start an MLflow run
with mlflow.start_run() as run:
    print("Starting MLflow run...")
    C_param = 0.75 # Using a different C value
    model = LogisticRegression(C=C_param, max_iter=200, random_state=42)
    model.fit(X_train, y_train)
    accuracy = accuracy_score(y_test, model.predict(X_test))

    # Log param and metric
    mlflow.log_param("C", C_param)
    mlflow.log_metric("accuracy", accuracy)

# --- THIS IS THE KEY STEP ---
# Log the model and register it under the specified name
mlflow.sklearn.log_model(
    model,
    "model",
    registered_model_name=REGISTERED_MODEL_NAME
)
# --- END KEY STEP ---

print(f"Run complete. Model registered as '{REGISTERED_MODEL_NAME}'")

```

5. Launch MLflow UI (in a separate terminal)

```

cd mlflow-practical
.\venv\Scripts\activate
mlflow ui

```

6. Run the registration script (in the first terminal)

```
python register_model.py
```

7. Check the MLflow UI

Go to <http://127.0.0.1:5000>. Click the "Models" tab. You will see "PracticalExamModel". Click it to see Version 1 and manage its stage (e.g., move to "Staging" or "Production").

❖ PRACTICAL 5: Automate using Prefect

GOAL: Create a simple ETL and training pipeline using Prefect. **NOTE:** Prefect is simpler for local setup than Airflow.

1. Setup Project Folder

```

mkdir prefect-practical
cd prefect-practical

```

2. Create and Activate Virtual Environment

```
python -m venv venv
.\venv\Scripts\activate
```

3. Install Libraries

```
pip install prefect scikit-learn pandas
```

4. Create Python file: pipeline.py

```
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from prefect import task, flow

@task
def extract_data():
    """Generates dummy data and saves it."""
    X, y = make_classification(n_samples=100, n_features=10, random_state=42)
    df = pd.DataFrame(X, columns=[f'feature_{i}' for i in range(10)])
    df['target'] = y
    df.to_csv('dummy_data.csv', index=False)
    print("Task: Data extracted and saved.")
    return 'dummy_data.csv'

@task
def transform_data(data_path: str):
    """Loads and splits the data."""
    df = pd.read_csv(data_path)
    X = df.drop('target', axis=1)
    y = df['target']
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
    print("Task: Data transformed.")
    return X_train, X_test, y_train, y_test

@task
def load_model(X_train, y_train, X_test, y_test):
    """Trains a model and prints accuracy."""
    model = LogisticRegression()
    model.fit(X_train, y_train)
    preds = model.predict(X_test)
    acc = accuracy_score(y_test, preds)
    print(f"Task: Model trained. Accuracy: {acc}")
    return acc

@flow(name="Practical Exam Flow")
def run_pipeline():
    """Main flow to run the ETL and training pipeline."""
    data_path = extract_data()
    X_train, X_test, y_train, y_test = transform_data(data_path)
    accuracy = load_model(X_train, y_train, X_test, y_test)
    print(f"Flow complete. Final accuracy: {accuracy}")

if __name__ == "__main__":
    run_pipeline()
```

5. Run the pipeline

```
python pipeline.py
```

6. (Optional) Start Prefect UI

```
# In a new terminal
prefect server start
```

Goto <http://127.0.0.1:4200> to see the dashboard. Run `python pipeline.py` again, and your flow run will appear in the UI.

¶ PRACTICAL 6: Build REST API (FastAPI) & Containerize (Docker)

GOAL: Train a model, wrap it in a FastAPI, and run it as a Docker container. **PRE-REQUISITE:** Docker Desktop must be installed AND RUNNING.

1. Setup Project Folder

```
mkdir fastapi-docker-practical
cd fastapi-docker-practical
```

2. Create and Activate Virtual Environment

```
python -m venv venv
.\venv\Scripts\activate
```

3. Install Libraries

```
pip install fastapi "uvicorn[standard]" scikit-learn pandas joblib
```

4. Create Python file: `create_model.py`

```
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.linear_model import LogisticRegression
import joblib

X, y = make_classification(n_samples=100, n_features=4, n_informative=2, n_redundant=0, random_state=42)
model = LogisticRegression()
model.fit(X, y)

joblib.dump(model, 'model.joblib')
print("Model saved to model.joblib")
```

5. Create Python file: `main.py`

```

from fastapi import FastAPI
from pydantic import BaseModel
import joblib
import numpy as np

app = FastAPI()

# Load the trained model
model = joblib.load('model.joblib')

# Define the input data model
class ModelInput(BaseModel):
    feature1: float
    feature2: float
    feature3: float
    feature4: float

@app.get("/")
def read_root():
    return {"message": "Model API is running."}

@app.post("/predict")
def predict(data: ModelInput):
    # Convert input data to numpy array
    input_data = np.array([[data.feature1, data.feature2, data.feature3, data.feature4]])

    # Get prediction and probability
    prediction = model.predict(input_data)[0]
    probability = model.predict_proba(input_data).max()

    return {
        "prediction": int(prediction),
        "probability": float(probability)
    }

```

6. Create file: requirements.txt

```

fastapi
uvicorn
scikit-learn
pandas
joblib
numpy

```

7. Create file: Dockerfile

```
# Use an official Python runtime as a parent image
FROM python:3.9-slim

# Set the working directory in the container
WORKDIR /app

# Copy the requirements file into the container
COPY requirements.txt .

# Install any needed packages specified in requirements.txt
RUN pip install --no-cache-dir -r requirements.txt

# Copy the rest of the application code
COPY . .

# Expose the port the app runs on
EXPOSE 8000

# Define the command to run the application
CMD ["uvicorn", "main:app", "--host", "0.0.0.0", "--port", "8000"]
```

8. Run the model training script

```
python create_model.py
# (This creates model.joblib)
```

9. Test locally (Optional)

```
uvicorn main:app --reload
```

Open <http://127.0.0.1:8000/docs> to test. Press Ctrl+C to stop.

10. Build the Docker image

```
docker build -t practical-api .
```

11. Run the Docker container

```
docker run -d -p 8000:8000 --name my-api practical-api
```

Test the container at <http://127.0.0.1:8000/docs>

12. Stop and remove the container

```
docker stop my-api
docker rm my-api
```

¶ PRACTICAL 7: Deploy with CI/CD (GitHub Actions)

GOAL: Automatically build the Docker image from Practical #6 when pushing to GitHub. **PRE-REQUISITES:**

- Project from Practical #6.
- A new, empty GitHub repository.
- Git installed.

1. Initialize Git in your project folder (from Practical #6)

```
cd fastapi-docker-practical
git init
git add .
git commit -m "Initial commit of FastAPI project"
```

2. Link to GitHub repository

```
# Get this command from your new GitHub repo page
git remote add origin https://github.com/YOUR_USERNAME/YOUR_REPO_NAME.git
git branch -M main
git push -u origin main
```

3. Create GitHub Actions workflow folder

```
mkdir .github
cd .github
mkdir workflows
cd workflows
```

4. Create file: `main.yml` (inside `.github/workflows`)

```
name: CI-CD Pipeline

on:
  push:
    branches: [ "main" ] # Trigger on push to main branch

jobs:
  build:
    runs-on: ubuntu-latest # Use a Linux runner

    steps:
      - name: Check out the repo
        uses: actions/checkout@v3

      - name: Set up QEMU
        uses: docker/setup-qemu-action@v2

      - name: Set up Docker Buildx
        uses: docker/setup-buildx-action@v2

      - name: Log in to GitHub Container Registry
        uses: docker/login-action@v2
        with:
          registry: ghcr.io
          username: ${{ github.repository_owner }}
          password: ${{ secrets.GITHUB_TOKEN }} # This is automatically provided

      - name: Build and push Docker image
        uses: docker/build-push-action@v4
        with:
          context: . # Use the root of our repo
          push: true
          tags: ghcr.io/${{ github.repository_owner }}/practical-api:latest # Tag image
```

5. Commit and push the workflow file

```
cd ../../ # Go back to the project root (fastapi-docker-practical)
git add .github/workflows/main.yml
git commit -m "Add GitHub Actions CI/CD workflow"
git push origin main
```

6. Check the "Actions" tab

Go to your GitHub repository in the browser. Click the "Actions" tab. You will see your workflow running. When it's done, check the "Packages" tab on your repo's main page to see your published Docker image.

¶ PRACTICAL 8: Monitor model with Grafana (and Prometheus)

GOAL: Expose metrics from the FastAPI (Practical #6) and view them in Grafana. **PRE-REQUISITES:** Docker Desktop (with Docker Compose) must be installed AND RUNNING.

1. Setup Project Folder (Use Practical #6)

```
cd fastapi-docker-practical  
.venv\Scripts\activate
```

2. Install Prometheus instrumentator

```
pip install prometheus-fastapi-instrumentator  
pip freeze > requirements.txt
```

3. Modify `main.py` to add the metrics endpoint

```
from fastapi import FastAPI  
from pydantic import BaseModel  
import joblib  
import numpy as np  
from prometheus_fastapi_instrumentator import Instrumentator # Import  
  
app = FastAPI()  
  
# --- ADD THIS BLOCK ---  
@app.on_event("startup")  
async def startup():  
    """Instrument the app on startup."""  
    Instrumentator().instrument(app).expose(app)  
# --- END BLOCK ---  
  
# Load the trained model  
model = joblib.load('model.joblib')  
  
# Define the input data model  
class ModelInput(BaseModel):  
    feature1: float  
    feature2: float  
    feature3: float  
    feature4: float  
  
@app.get("/")  
def read_root():  
    return {"message": "Model API is running."}  
  
@app.post("/predict")  
def predict(data: ModelInput):  
    input_data = np.array([[data.feature1, data.feature2, data.feature3, data.feature4]])  
    prediction = model.predict(input_data)[0]  
    probability = model.predict_proba(input_data).max()  
    return {  
        "prediction": int(prediction),  
        "probability": float(probability)  
    }
```

4. Modify Dockerfile (No change needed if `requirements.txt` was updated)

Just rebuild the image to include the new library.

```
docker build -t practical-api .
```

5. Create file: `prometheus.yml` (for Prometheus config)

```

global:
  scrape_interval: 15s

scrape_configs:
  - job_name: 'fastapi-app'
    static_configs:
      # Use 'host.docker.internal' to allow Docker container
      # to access the API running on the host machine.
      - targets: ['host.docker.internal:8000']

```

6. Create file: docker-compose.yml (to run Grafana/Prometheus)

```

version: '3.8'

services:
  prometheus:
    image: prom/prometheus:latest
    container_name: prometheus
    ports:
      - "9090:9090"
    volumes:
      - ./prometheus.yml:/etc/prometheus/prometheus.yml
    command:
      - '--config.file=/etc/prometheus/prometheus.yml'

  grafana:
    image: grafana/grafana-oss:latest
    container_name: grafana
    ports:
      - "3000:3000"
    depends_on:
      - prometheus
    environment:
      - GF_SECURITY_ADMIN_USER=admin
      - GF_SECURITY_ADMIN_PASSWORD=admin

```

7. Run everything

```

# Terminal 1: Run Docker Compose (Prometheus + Grafana)
docker-compose up

# Terminal 2: Run the FastAPI app (on the host)
.\venv\Scripts\activate
uvicorn main:app --host 0.0.0.0 --port 8000

```

8. Configure Grafana

- Open Grafana: <http://127.0.0.1:3000> (user: admin, pass: admin)
- Add Data Source:
 - Connections > Add new connection > Prometheus
 - URL: <http://prometheus:9090>
 - Click "Save & Test".
- Create Dashboard:
 - Dashboards > New > Add visualization
 - Select "Prometheus" as data source.
 - In the "Metric" browser, try: `http_requests_total`
 - Click "Run queries". You will see graphs.
- Simulate Drift (Optional): Send traffic to your API (e.g., using `curl` or Postman) to see the graphs change.

9. Stop everything

```

# In Terminal 1
docker-compose down

```

GOAL: Train a model and use SHAP to explain its predictions.

1. Setup Project Folder

```
mkdir shap-practical  
cd shap-practical
```

2. Create and Activate Virtual Environment

```
python -m venv venv  
.\\venv\\Scripts\\activate
```

3. Install Libraries

```
pip install shap scikit-learn pandas matplotlib
```

4. Create Python file: `create_data.py`

```
import pandas as pd  
from sklearn.datasets import make_classification  
  
X, y = make_classification(n_samples=100, n_features=10, n_informative=5, n_redundant=0, random_state=42)  
df = pd.DataFrame(X, columns=[f'feature_{i}' for i in range(10)])  
df['target'] = y  
  
df.to_csv('dummy_data.csv', index=False)  
print("Data created.")
```

5. Create Python file: `explain.py`

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
import shap
import matplotlib.pyplot as plt

# Load data
df = pd.read_csv('dummy_data.csv')
X = df.drop('target', axis=1)
y = df['target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train a model
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)
print(f"Model trained. Accuracy: {model.score(X_test, y_test)}")

# Initialize SHAP Explainer
# For tree models, use shap.TreeExplainer for speed
explainer = shap.TreeExplainer(model)

# Calculate SHAP values for the test set
shap_values = explainer.shap_values(X_test)

# Generate summary plot (beeswarm)
# shap_values[1] is for the "positive" class (class 1)
shap.summary_plot(shap_values[1], X_test, show=False)

# Save the plot to a file
plt.savefig('shap_summary.png', bbox_inches='tight')
print("SHAP summary plot saved to shap_summary.png")

# Generate plot for a single prediction
plt.clf() # Clear the previous plot
shap.force_plot(explainer.expected_value[1], shap_values[1][0,:], X_test.iloc[0,:], matplotlib=True, show=False)
plt.savefig('shap_force_plot.png', bbox_inches='tight')
print("SHAP force plot saved to shap_force_plot.png")

```

6. Run setup and explanation

```

python create_data.py
python explain.py

```

Check the folder for `shap_summary.png` and `shap_force_plot.png`